


Threads & Tasks: Executor Framework

- **Introduction & Motivation**
 - WebServer
- **Executor Framework**
- **Callable and Future**
- **Fork-Join Motivation**
- **Fork-Join Framework**

Threads & Tasks

- Example: Webserver



Introduction: Implementation of a WebServer1

```
public class WebServer1 {  
  
    public static void main(String[] args) throws IOException {  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true) {  
            Socket s = serverSocket.accept();  
            handleRequest(s);  
        }  
    }  
}
```

- **Single-Threaded Server:**
Tasks are executed within the server thread
 - Poor performance, can only handle one request at a time
 - While server is processing, new requests must wait
 - Possible if request processing is very fast

Introduction: Implementation of a WebServer2

```
public class WebServer2 {  
  
    public static void main(String[] args) throws IOException {  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true) {  
            Socket s = serverSocket.accept();  
            Thread t = new Thread(() -> handleRequest(s));  
            t.start();  
        }  
    }  
}
```

- **Explicitly creating threads for tasks**
Each task is executed in its own thread
 - handleRequest must be thread-safe
 - Excessive thread creation: Scheduling overhead / memory consumption

Introduction: Implementation of a WebServer3

```
public class WebServer3 {  
  
    public static void main(String[] args) throws IOException {  
        final ServerSocket serverSocket = new ServerSocket(80);  
        for (int i = 0; i < 10; i++) {  
            Thread t = new Thread(() -> {  
                while (true) {  
                    try { handleRequest(serverSocket.accept()); }  
                    catch (IOException e) { /* ... */ }  
                }  
            });  
            t.start();  
        }  
    }  
}
```

Introduction: Implementation of a WebServer3

- **Disadvantages**

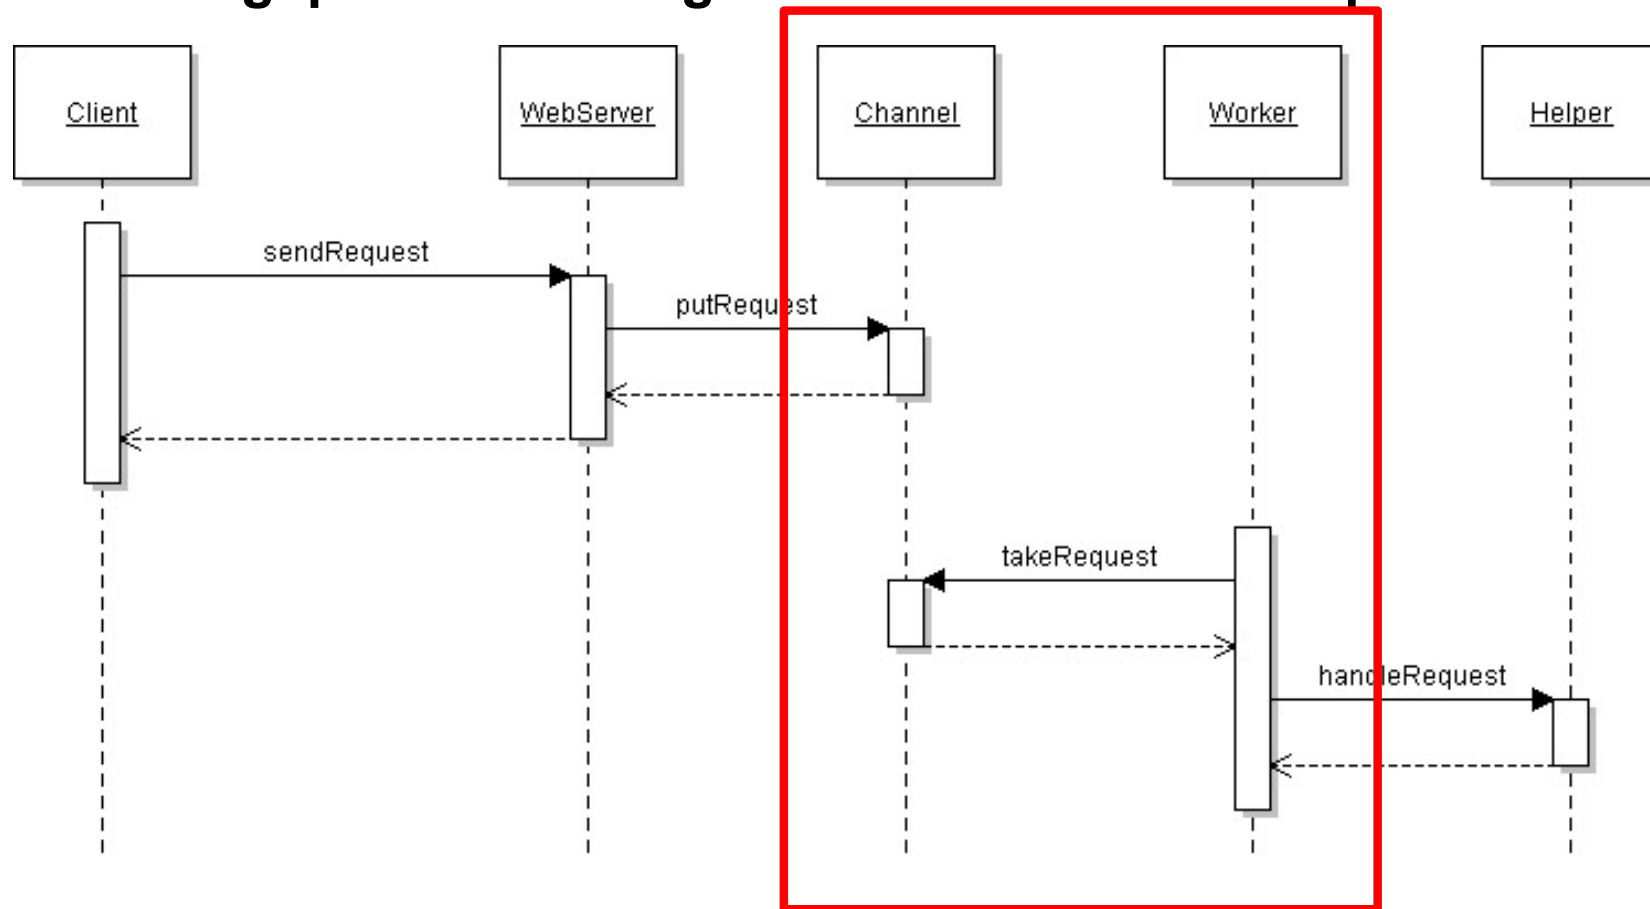
- Maybe incorrect, `ServerSocket.accept()` is *not* documented as threadsafe
- No flexibility, i.e. exactly 10 threads
 - No creation of new threads
 - No deletion of unused threads
- All threads are pre-created, no lazy allocation
- No life-cycle management, i.e. "pool" cannot be stopped
- No error handling, if an exception is not caught, thread terminates silently
- No flexibility in the order of pending requests
 - They are stored in the Server-Socket queue
- No full control over queue for pending requests
 - backlog parameter only specifies *maximum* length of the queue

Threads & Tasks: Executor Framework

- Introduction & Motivation
- **Executor Framework**
- Callable and Future
- Fork-Join Motivation
- Fork-Join Framework

Executor Framework

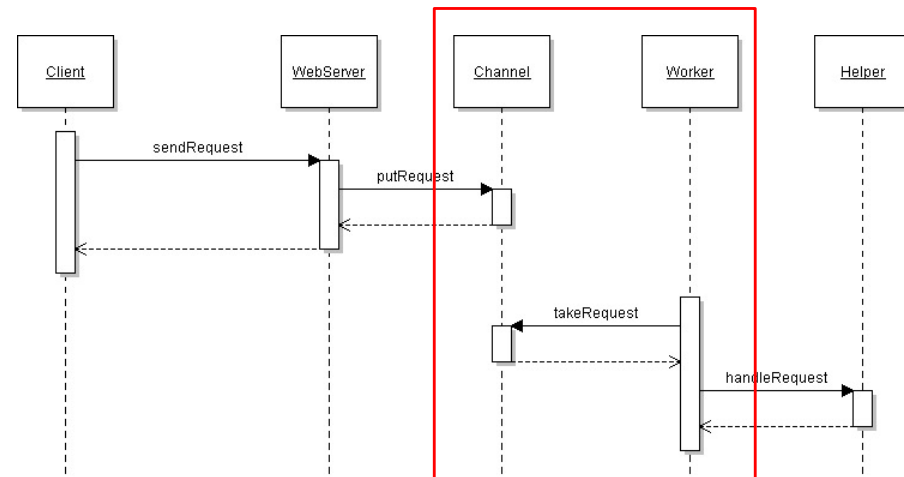
- Fills the gap between single-threaded and thread-per-task



Executor Framework

- **Participants**

- Worker: One thread is used to execute many unrelated tasks
 - These threads are called worker threads / background threads
 - May be organized in a thread pool (if more than one thread is used)
 - May provide a flexible thread management
- Channel: A buffer which holds pending requests
 - May be bounded
 - Implements the producer-consumer pattern



Executor Framework

- **Executor = Channel + Worker**

```
public interface Executor {  
    void execute(Runnable task);  
}
```

- Decouples task submission from task execution
- Executes the given task at some time in the future
- The task may be executed
 - in a new thread / in a pooled thread / in the calling thread

- **Runnable = Task**

```
public interface Runnable {  
    void run();  
}
```

- Limitation:
 - Method *run* cannot return a result (results are placed in shared fields)
 - Method *run* cannot declare a checked exception

Introduction: Implementation of a WebServer4

```
public class WebServer4 {  
    public static void main(String[] args) throws IOException {  
        Executor exec = new ThreadPoolExecutor(10);  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true){  
            final Socket s = serverSocket.accept();  
            Runnable task = new Runnable(){  
                public void run(){ handleRequest(s); }  
            };  
            exec.execute(task);  
        }  
    }  
}
```

- Using lambda notation

```
exec.execute(() -> handleRequest(s));
```

Executor: Implementation

```
class MyThreadPoolExecutor implements Executor {  
    private final BlockingQueue<Runnable> queue  
        = new LinkedBlockingQueue<Runnable>();  
  
    public void execute(Runnable r) { queue.offer(r); }  
  
    public MyThreadPoolExecutor(int nrThreads) {  
        for (int i = 0; i < nrThreads; i++) { activate(); }  
    }  
  
    private void activate() {  
        new Thread(() -> {  
            try {  
                while (true) { queue.take().run(); }  
            } catch (InterruptedException e) { /* die */ }  
        }).start();  
    }  
}
```

Executor: Simple implementations

- **DirectExecutor: Synchronous execution (in calling thread)**

```
class DirectExecutor implements Executor {  
    public void execute(Runnable r) { r.run(); }  
}
```

- With this executor Webserver4 = Webserver1

- **Thread per task executor**

```
class ThreadPerTaskExecutor implements Executor {  
    public void execute(Runnable r) {  
        new Thread(r).start();  
    }  
}
```

- With this executor Webserver4 = Webserver2

Executor: Advanced Implementations

- **Execution Policies**

- Execution order of submitted tasks (FIFO, LIFO, Priority Queue)
- Number of threads which execute concurrently
- Maximal size of queue with pending tasks
- Actions taken before / after task execution
 - E.g. extension hooks in class `ThreadPoolExecutor`
 - `void beforeExecute(Thread t, Runnable r) { }`
 - `void afterExecute(Runnable r, Throwable t) { }`

- **Factory: `java.util.concurrent.Executors`**

- Factory methods for preconfigured `ThreadPoolExecutor` instances
- E.g. creating an executor using 10 worker threads:

```
Executor es = Executors.newFixedThreadPool(10);
```

Executor: Advanced Implementations

- **Executors.newFixedThreadPool**
 - Threads are created up to a fixed number
 - Threads which die due to an unexpected exception are replaced
- **Executors.newCachedThreadPool**
 - Creates new threads as needed, reusing previously constructed threads if they are available
- **Executors.newSingleThreadExecutor**
 - Uses single worker thread
 - Worker thread is replaced if an unexpected exception occurs
- **Executors.newScheduledThreadPool**
 - Creates a ScheduledExecutorService which supports
 - Periodic tasks (scheduleAtFixedRate / scheduleWithFixedDelay)
 - Delayed tasks (schedule)

ThreadFactory

- **Some factory methods on j.u.c.Executors take a ThreadFactory**

```
public interface ThreadFactory {  
    Thread newThread(Runnable r);  
}
```

- **Enables applications to use**
 - special Thread subclasses
 - custom named threads
 - daemon flag

Executors: Design Considerations

- **Identity**

- Different tasks are executed with the same thread
 - If task uses ThreadLocals it needs to clean up afterwards
- Logical flows can be executed on multiple threads
 - Information in a ThreadLocal is not available inside the whole flow

➤ There is no more a one-to-one relation between a logical flow and a technical Thread

- **Granularity**

- Promotes concurrency by providing a natural structure for parallelizing work

ExecutorService: Executor Life-Cycle

- **ExecutorService: provides life-cycle management methods**

```
interface ExecutorService extends Executor {  
    void shutdown();  
    List<Runnable> shutdownNow();  
    boolean isShutdown();  
    boolean isTerminated();  
    boolean awaitTermination(long timeout, TimeUnit unit)  
        throws InterruptedException;  
}
```

- Supports shutdown methods
 - shutdown: Graceful shutdown: finish pending tasks, do not accept new ones
 - shutdownNow: Abrupt shutdown: running tasks are interrupted, returns list of tasks that were not started
- awaitTermination: awaits until executor service is terminated



Executors and JMM

- **Memory consistency effects**
 - Actions in a thread prior to submitting a Runnable object to an **Executor** **happen-before** its execution begins (possibly in another thread)
 - Actions in a task which is executed by a **SingleThreadExecutor** **happen-before** actions executed in subsequent tasks (even if the subsequent task is executed by another thread due to an exception)

Threads & Tasks: Executor Framework

- Introduction & Motivation
- Executor Framework
- **Callable and Future**
- Fork-Join Motivation
- Fork-Join Framework

Result-bearing tasks: Callable & Future

- **Callable: Task with a result / exception**

```
interface Callable<V> {  
    V call() throws Exception;  
}
```

- **Future: represents a future result of a task**

```
interface Future<V> {  
    boolean cancel(boolean mayInterruptIfRunning);  
    boolean isCancelled();  
    boolean isDone();  
    V get() throws InterruptedException, ExecutionException,  
        CancellationException;  
    V get(long timeout, TimeUnit unit)  
        throws InterruptedException, ExecutionException,  
        CancellationException, TimeoutException;  
}
```


Callable & Future

- **Submitting tasks**

```
interface ExecutorService extends Executor {  
    // ... Lifecycle methods  
    <T> Future<T> submit(Callable<T> task);  
  
    Future<?> submit(Runnable task);  
    <T> Future<T> submit(Runnable task, T result);  
  
    <T> List<Future<T>> invokeAll(  
        Collection<? extends Callable<T>> tasks)  
        throws InterruptedException;  
  
    <T> T invokeAny(  
        Collection<? extends Callable<T>> tasks)  
        throws InterruptedException, ExecutionException;  
}
```

Callable & Future

- **States of a Task**



- **get()**
 - If completed: returns immediately
(returns result or throws `ExecutionException`)
 - If not completed: method call blocks
 - If terminates regularly => result
 - If terminates with an exception => `ExecutionException`
 - If cancelled => `CancellationException`
 - If thread calling get was interrupted => `InterruptedException`

CompletionService

- **CompletionService = Executor + BlockingQueue**
 - Decouples production of new tasks from the consumption of the results of completed tasks
 - Producers submit tasks for execution
 - Consumers take completed tasks and process their results

```
interface CompletionService {  
    Future<V> submit(Callable<V> task);  
    Future<V> submit(Runnable task, V result);  
    Future<V> take() throws IE; // waits for a future  
    Future<V> poll(); // returns available future or null  
    Future<V> poll(long timeout, TimeUnit unit) throws IE;  
}
```

- take retrieves and removes the next completed task, potentially waiting

- **ExecutorCompletionService**
 - Uses a separate Executor to schedule the tasks

Callable & Future and JMM

- **Memory consistency effects**
 - Actions in a thread prior to the submission of a Runnable or Callable task to an ExecutorService **happen-before** any actions taken by that task
 - Any actions taken by a Runnable or Callable task executed by an ExecutorService **happen-before** the result is retrieved via Future.get()

Threads & Tasks: Executor Framework

- Introduction & Motivation
- Executor Framework
- Callable and Future
- **Fork-Join Motivation**
- Fork-Join Framework

MergeSort using Executor

```
public class MS1 implements Runnable {
    public final int[] is, tmp; private final int l, r;
    private final ExecutorService ex;

    public MS1(int[] is, int[] tmp, int l, int r, ExecutorService ex) {
        this.is = is; this.tmp = tmp; this.l = l; this.r = r; this.ex = ex;
    }

    public void run() {
        if(r - l <= 1) return;
        else {
            int mid = (l+r) / 2;
            MS1 left = new MS1(is, tmp, l, mid, ex);
            MS1 right = new MS1(is, tmp, mid, r, ex);

            Future<?> lf = ex.submit(left);
            Future<?> rf = ex.submit(right);

            try { lf.get(); rf.get(); } catch (InterruptedException | ExecutionException e) {}
            merge(is, tmp, l, mid, r);
        }
    }

    private void merge(int[] is, int[] tmp, int l, int m, int r) {...}
}
```

MergeSort using Executor

```
int[] data = ...
int[] tmp = new int[data.length];
ExecutorService es = Executors.newFixedThreadPool(3);

MS1 ms = new MS1(data, tmp, 0, data.length-1, es);
Future<?> f = es.submit(ms);
f.get();
// When this line is reached, ms.is is sorted!
```

MergeSort using hard-working Executor

```
public class MS2 implements Runnable {
    ...
    public void run() {
        if(r - l <= 1) return;
        else {
            int mid = (l+r) / 2;
            MS2 left = new MS2(is, tmp, l, mid, hex);
            MS2 right = new MS2(is, tmp, mid, r, hex);
            Future<?> lf = ex.submit(left);
            Future<?> rf = ex.submit(right);

            // don't hang around, work!
            while (!(lf.isDone() && rf.isDone())) {
                ex.helpOtherWorkers();
            }
            try {lf.get(); rf.get();} catch (InterruptedException | ExecutionException e) {}
            merge(is, tmp, l, mid, r);
        }
    }
}
```

MergeSort using hard-working Executor

```
public class MS2 implements Runnable {
    ...
    public void run() {
        if(r - l <= 1000) Arrays.sort(is); return;
        else {
            int mid = (l+r) / 2;
            MS2 left = new MS2(is, tmp, l, mid, hex);
            MS2 right = new MS2(is, tmp, mid, r, hex);
            Future<?> lf = ex.submit(left);
            Future<?> rf = ex.submit(right);

            // don't hang around, work!
            while (!(lf.isDone() && rf.isDone())) {
                ex.helpOtherWorkers();
            }
            try {lf.get(); rf.get();} catch (InterruptedException | ExecutionException e) {}
            merge(is, tmp, l, mid, r);
        }
    }
}
```

ForkJoin decomposition



- **Parallel version of divide-and-conquer algorithms**

```
Result solve(Problem problem) {  
    if (problem.size < SEQUENTIAL_THRESHOLD)  
        return solveSequentially(problem);  
    else {  
        Result left, right;  
        INVOKE-PARALLEL {  
            left = solve(extractLeftHalf(problem));  
            right = solve(extractRightHalf(problem));  
        }  
        return combine(left, right);  
    }  
}
```

Examples of divide-and-conquer algorithms

- **Sort**
 - Mergesort
 - Quicksort
- **Search in unsorted data**
 - Arrays
 - Trees
- **Numerics**
 - Matrix multiplications
- **Geometrics**
 - Convex hull


Threads & Tasks: Executor Framework

- Introduction & Motivation
- Executor Framework
- Callable and Future
- Fork-Join Motivation
- **Fork-Join Framework**

ForkJoin Framework: Work stealing



- Create a limited number of worker threads
- Each worker thread maintains a private double-ended work queue
- When forking, worker pushes new task at the head of its deque
- When waiting or idle, worker pops a task off the head of its deque and executes it instead of sleeping
- If a worker's deque is empty, it steals a task off the tail of another randomly chosen worker



Work stealing: Efficiency

- **Reduced contention compared to shared work queue**
 - No contention for head since it is only accessed by the owner
 - No contention between head and tail access (good deque algorithms)
 - Almost never contention for tail because stealing is infrequent
 - Workers access their deque in LIFO (Stack) order
 - Size of tasks gets smaller as problem is divided
 - Tasks are stolen from the tail and are generally a big chunk!
- **Good load balancing**
 - If work is unequally distributed, it is rebalanced via stealing
 - Without central coordination
 - With little scheduling overhead
 - With minimal synchronization costs

Task granularity and structure

- **Small Tasks (low threshold)**
 - Maximizing parallelism
 - More fine-grained tasks keep more CPUs busy
 - Improves load balancing and locality
 - Decreases time that CPUs must wait for one another
 - Fast if data fits into cache
- **Large Tasks (high threshold)**
 - Minimizing overhead
 - Task creation and management versus sequential execution
 - Memory consumption, garbage collection

Fork-join task framework is designed to minimize per task overhead for compute-intensive tasks.

MergeSort using ForkJoin

```
public class FJMS extends RecursiveAction {
    public final int[] is, tmp; private final int l, r;

    public FJMS(int[] is, int[] tmp, int l, int r) {
        this.is = is; this.tmp = tmp; this.l = l; this.r = r;
    }

    protected void compute() {
        if (r - l <= 100000) Arrays.sort(elems, l, r);
        else {
            int mid = (l + r) / 2;
            FJMS left = new FJMS (is, tmp, l, mid);
            FJMS right = new FJMS (is, tmp, mid, r);
            left.fork();
            right.invoke();
            left.join();
            merge(is, tmp, l, mid, r);
        }
    }


    private void merge(int[] es, int[] tmp, int l, int m, int r) {...}
}
```

Executing FJ Tasks: ForkJoinPool

```
public class ForkJoinRunner {  
    public static void main(String[] args) {  
        int[] data = ...  
        int[] tmp = new int[data.length];  
        ForkJoinPool fjPool = new ForkJoinPool();  
        FJMS ms = new FJMS(data,tmp,0,data.length);  
        fjPool.invoke(ms);  
        // When this line is reached, ms.is is sorted!  
    }  
}
```


- Default constructor of ForkJoinPool sets parallelism to `Runtime.getRuntime().availableProcessors()`
- Normally a single ForkJoinPool is used for all parallel task execution in a program or subsystem: **`ForkJoinPool.commonPool()`**
- Because it uses daemon threads, there is no need for explicit shutdown

ForkJoin Framework: Overview



- A ForkJoinPool maintains typically #CPU ForkJoinWorkerThreads
- Each ForkJoinWorkerThread executes several ForkJoinTasks
- Your tasks subclass ForkJoinTask

RecursiveAction vs. RecursiveTask



`protected abstract void compute();`

- Resultless
- Typically reuses argument field for result

`protected abstract V compute();`

- Result-bearing

ForkJoinPool: API

```
public class ForkJoinPool extends AbstractExecutorService {  
    // ... ExecutorService methods  
  
    // Constructors  
    public ForkJoinPool(int parallelism) { ... }  
  
    // Task processing  
    public <T> T invoke(ForkJoinTask<T> task) { ... } //block.  
    public void execute(ForkJoinTask<?> task) { ... } //nonblock.  
    public <T> ForkJoinTask<T> submit(ForkJoinTask<T> task) { ... }  
  
    // Monitoring operations (used for tuning)  
    public long getStealCount() { ... }  
    public long getQueuedTaskCount() { ... }  
    public int getQueuedSubmissionCount() { ... }  
  
    ...  
}
```

ForkJoinTask: API

```
public abstract class ForkJoinTask<V> implements Future<V>, Serializable
{
    // ... Future methods

    // Subtask control
    public final ForkJoinTask<V> fork() { ... } // put into deque
    public final V join() { ... } // work on other tasks
    public final V invoke() { ... } // invoke with current FJWT

    // Bulk invocations
    public static void invokeAll(ForkJoinTask<?> t1,
                                ForkJoinTask<?> t2) { ... }
    public static void invokeAll(ForkJoinTask<?>... tasks) { ... }
    public static <T extends ForkJoinTask<?>> Collection<T>
        invokeAll(Collection<T> tasks) { ... }

    // adapters
    public static ForkJoinTask<?> adapt(Runnable runnable) { ... }
    ...
}
```

Summary: ForkJoin Framework

- **Framework for executing computational tasks**
 - Offers a portable way to express many parallel algorithms
 - Optimized for fine grained tasks
 - Based on work stealing
 - Ready for the massive multi core future
- **Usage**
 - Implement a subclass of RecursiveAction or RecursiveTask
 - Execute it on a ForkJoinPool